

# deepdraw

February 24, 2017

```
In [ ]: # imports and basic notebook setup
from cStringIO import StringIO
import numpy as np
import os, re, random
import scipy.ndimage as nd
import PIL.Image
import sys
from IPython.display import clear_output, Image, display
from scipy.misc import imresize
import cv2
import matplotlib.pyplot as plt

%matplotlib inline

pycaffe_root = "/home/s/caffe/python"
sys.path.insert(0, pycaffe_root)
import caffe

model_name = "GoogLeNet"
model_path = '/home/s/caffe/models/bvlc_googlenet/'
net_fn = './deploy_googlenet_updated.prototxt'
param_fn = model_path + 'bvlc_googlenet.caffemodel'
mean = np.float32([104.0, 117.0, 123.0])

net = caffe.Classifier(net_fn, param_fn,
                      mean = mean, # ImageNet mean, training set dependent
                      channel_swap = (2,1,0)) # the reference model has channels in BGR order

# a couple of utility functions for converting to and from Caffe's input image layout
def preprocess(net, img):
    return np.float32(np.rollaxis(img, 2)[::-1]) - net.transformer.mean['data']
def deprocess(net, img):
    return np.dstack((img + net.transformer.mean['data'])[::-1])

def blur(img, sigma):
    if sigma > 0:
        img[0] = nd.filters.gaussian_filter(img[0], sigma, order=0)
        img[1] = nd.filters.gaussian_filter(img[1], sigma, order=0)
        img[2] = nd.filters.gaussian_filter(img[2], sigma, order=0)
    return img

def bilateralFilter_and_Blur(img, sigma):
    img[0] = cv2.bilateralFilter(nd.filters.gaussian_filter(img[0], sigma, order=0), -1, 5, 2)
```

```

img[1] = cv2.bilateralFilter(nd.filters.gaussian_filter(img[1], sigma, order=0), -1, 5, 2)
img[2] = cv2.bilateralFilter(nd.filters.gaussian_filter(img[2], sigma, order=0), -1, 5, 2)
return img

def no_blur(img, sigma):
    # do nothing
    return img

def showarray(a, f, fmt='jpeg'):
    a = np.uint8(np.clip(a, 0, 255))
    f = StringIO()
    PIL.Image.fromarray(a).save(f, fmt)
    display(Image(data=f.getvalue()))

```

Definition of the main gradient ascent functions. Based on the [deeptream code](#) published by Google.

```

In [ ]: def make_step(net, step_size=1.5, end='inception_4c/output', clip=True, focus=None, sigma=None)
        '''Basic gradient ascent step.'''

        src = net.blobs['data'] # input image is stored in Net's 'data' blob

        dst = net.blobs[end]
        net.forward(end=end)

        one_hot = np.zeros_like(dst.data)
        one_hot.flat[focus] = 1.
        dst.diff[:] = one_hot

        net.backward(start=end)
        g = src.diff[0]

        src.data[:] += step_size/np.abs(g).mean() * g

        if clip:
            bias = net.transformer.mean['data']
            src.data[:] = np.clip(src.data, -bias, 255-bias)

        # Note: All of these
        src.data[0] = bilateralFilter_and_Blur(src.data[0], sigma) # Approach #3 (blur + bilateral)
        # src.data[0] = blur(src.data[0], sigma) # Approach #3 (only blur)
        # src.data[0] = no_blur(src.data[0], sigma) # Approach #2

        # reset objective for next step
        dst.diff.fill(0.)

def deepdraw(net, base_img, octaves, random_crop=True, focus=None,
            clip=True, **step_params):

    # prepare base image
    image = preprocess(net, base_img)

    # get input dimensions from net
    w = net.blobs['data'].width
    h = net.blobs['data'].height

```

```

print "starting drawing"
src = net.blobs['data']
src.reshape(1,3,h,w) # reshape the network's input image
for e,o in enumerate(octaves): # scale up image if specified in oc
    if 'scale' in o:
        # resize by o['scale'] if it exists
        image = nd.zoom(image, (1,o['scale'],o['scale']))
    _,imw,imh = image.shape

    # select layer
    layer = o['layer']

    for i in xrange(o['iter_n']):
        if imw > w:
            if random_crop:
                # randomly select a crop
                mid_x = (imw-w)/2.
                width_x = imw-w
                ox = np.random.normal(mid_x, width_x*0.3, 1)
                ox = int(np.clip(ox,0,imw-w))
                mid_y = (imh-h)/2.
                width_y = imh-h
                oy = np.random.normal(mid_y, width_y*0.3, 1)
                oy = int(np.clip(oy,0,imh-h))
                # insert the crop into src.data[0]
                src.data[0] = image[:,ox:ox+w,oy:oy+h]
            else:
                ox = (imw-w)/2.
                oy = (imh-h)/2.
                src.data[0] = image[:,ox:ox+w,oy:oy+h]
        else:
            ox = 0
            oy = 0
            src.data[0] = image.copy()

        sigma = o['start_sigma'] + ((o['end_sigma'] - o['start_sigma']) * i) / o['iter_n']
        step_size = o['start_step_size'] + ((o['end_step_size'] - o['start_step_size']) * i)

        make_step(net, end=layer, clip=clip, focus=focus,
                  sigma=sigma, step_size=step_size)

        if i % 10 == 0:
            print 'finished step %d in octave %d' % (i,e)

        # insert modified image back into original image (if necessary)
        image[:,ox:ox+w,oy:oy+h] = src.data[0]

    print "octave %d image:" % e
    showarray(deprocess(net, image), "./octave_"+str(e)+".jpg")

    # returning the resulting image
    return cv2.bilateralFilter((np.uint8(deprocess(net, image))), -1, 20, 8)

```

Generating the class visualizations

- `layer` : which layer to optimize
- `iter_n` : how many iterations
- `scale` : by what factor (if any) to scale up the base image before proceeding
- `start_sigma` : the initial radius of the gaussian blur
- `end_sigma` : the final radius of the gaussian blur
- `start_step_size` : the initial step size of the gradient ascent
- `end_step_size` : the final step size of the gradient ascent

```
In [ ]: # these octaves are where bulk of tuning occurs, they determine gradient ascent steps
octaves = [
    {
        'layer': 'loss3/classifier',    # layer to perform image updates from
        'iter_n': 190,                  # number of times to perform image update
        'start_sigma': 2.5,              # Gradually reduce gaussian blur as recommended by
        'end_sigma': 0.78,               # https://github.com/kylemcDonald/deepdream
        'start_step_size': 11.,          # Gradually change step size of gradient ascent as recomm
        'end_step_size': 11.,           # http://www.auduno.com/2015/07/29/visualizing-googlenc

    },
    {
        'layer': 'loss3/classifier',
        'scale': 1.2,
        'iter_n': 150,
        'start_sigma': 0.78*1.2,
        'end_sigma': 0.78,
        'start_step_size': 6.,
        'end_step_size': 6.

    },
    {
        'layer': 'loss3/classifier',
        'scale': 1.2,
        'iter_n': 70,
        'start_sigma': 0.78*1.2,
        'end_sigma': 0.44,
        'start_step_size': 6.,
        'end_step_size': 3.

    },
    {
        'layer': 'loss3/classifier',
        'scale': 1.2,
        'iter_n': 50,
        'start_sigma': 0.44,
        'end_sigma': 0.304,
        'start_step_size': 3.,
        'end_step_size': 3.

    },
    {
        'layer': 'loss3/classifier',
        'scale': 1.2,
        'iter_n': 30,
        'start_sigma': 0.44,
        'end_sigma': 0.304,
        'start_step_size': 3.,
        'end_step_size': 3.

    },
]
```

```
{
    'layer': 'loss2/classifier',
    'scale': 1.2,
    'iter_n': 150,
    'start_sigma': 0.44,
    'end_sigma': 0.304,
    'start_step_size': 3.,
    'end_step_size': 3.
},
{
    'layer': 'loss2/classifier',
    'scale': 1.2,
    'iter_n': 70,
    'start_sigma': 0.44,
    'end_sigma': 0.304,
    'start_step_size': 3.,
    'end_step_size': 3.
},
{
    'layer': 'loss2/classifier',
    'scale': 1.2,
    'iter_n': 40,
    'start_sigma': 0.44,
    'end_sigma': 0.304,
    'start_step_size': 3.,
    'end_step_size': 3.
},
{
    'layer': 'loss2/classifier',
    'scale': 1.2,
    'iter_n': 30,
    'start_sigma': 0.44,
    'end_sigma': 0.304,
    'start_step_size': 3.,
    'end_step_size': 3.
},
{
    'layer': 'loss2/classifier',
    'scale': 1.2,
    'iter_n': 20,
    'start_sigma': 0.44,
    'end_sigma': 0.304,
    'start_step_size': 3.,
    'end_step_size': 3.
},
{
    'layer': 'loss2/classifier',
    'scale': 1.2,
    'iter_n': 20,
    'start_sigma': 0.44,
    'end_sigma': 0.304,
    'start_step_size': 3.,
    'end_step_size': 3.
},
}
```

```

{
    'layer':'loss1/classifier',
    'scale':1.2,
    'iter_n':20,
    'start_sigma':0.44,
    'end_sigma':0.304,
    'start_step_size':3.,
    'end_step_size':3.
},
{
    'layer':'loss1/classifier',
    'scale':1.2,
    'iter_n':20,
    'start_sigma':0.44,
    'end_sigma':0.304,
    'start_step_size':3.,
    'end_step_size':3.
},
{
    'layer':'loss1/classifier',
    'iter_n':30,
    'start_sigma':0.44,
    'end_sigma':0.304,
    'start_step_size':3.,
    'end_step_size':3.
}
]

# get original image input size from network
original_w = net.blobs['data'].width
original_h = net.blobs['data'].height
# the background color of the initial image
background_color = np.float32([28.0, 84.0, 122.0]) # This is blue

for i in range(0,100): # Automate image visualization and writing to disk
    imagenet_class = i
    gen_image = np.random.normal(background_color, 8, (original_w, original_h, 3)) # create image
    gen_image = deepdraw(net, gen_image, octaves, focus=imagenet_class, random_crop=True)
    img_fn = '_' .join([model_name, "deepdraw_denoised", str(imagenet_class)+'.png'])
    PIL.Image.fromarray(cv2.bilateralFilter(np.uint8(gen_image), -1, 25, 6)).save('./' + img_fn)
    # PIL.Image.fromarray(np.uint8(gen_image)).save('./' + img_fn)

```

Less scaling and more iterations gives lower resolution but more coherent image

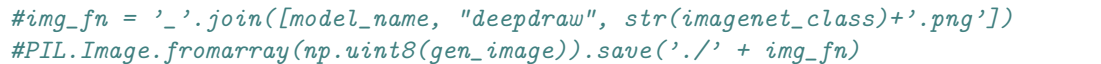
```

In [ ]: octaves = [
    {
        'layer':'loss3/classifier',
        'iter_n':190,
        'start_sigma':2.5,
        'end_sigma':0.78,
        'start_step_size':11.,
        'end_step_size':11.
    },
    {
        'layer':'loss3/classifier',

```

```

        'scale':1.2,
        'iter_n':450,
        'start_sigma':0.78*1.2,
        'end_sigma':0.40,
        'start_step_size':6.,
        'end_step_size':3.
    }
]
imagenet_class = 63
gen_image = np.random.normal(background_color, 8, (original_w, original_h, 3))
gen_image = deepdraw(net, gen_image, octaves, focus=imagenet_class,
                    random_crop=True)


#img_fn = '_' .join([model_name, "deepdraw", str(imagenet_class)+'.png'])
#PIL.Image.fromarray(np.uint8(gen_image)).save('./' + img_fn)

```

The choices below give images that are not as large as in first set of automation, but are more coherent

```

In [ ]: octaves = [
    {
        'layer':'loss3/classifier',

        'iter_n':190,
        'start_sigma':2.5,
        'end_sigma':0.78,
        'start_step_size':11.,
        'end_step_size':11.
    },
    {
        'layer':'loss3/classifier',
        'scale':1.2,
        'iter_n':100,
        'start_sigma':0.78*1.2,
        'end_sigma':0.65,
        'start_step_size':6.,
        'end_step_size':6.
    },
    {
        'layer':'loss2/classifier',
        'scale':1.2,
        'iter_n':90,
        'start_sigma':0.78*1,
        'end_sigma':0.55,
        'start_step_size':6.,
        'end_step_size':3.
    },
    {
        'layer':'loss2/classifier',
        'scale':1.2,
        'iter_n':60,
        'start_sigma':0.78*0.8,
        'end_sigma':0.45,
        'start_step_size':6.,
        'end_step_size':3.
    },
]

```

```

    {
        'layer':'loss1/classifier',
        'scale':1.2,
        'iter_n':10,
        'start_sigma':0.45,
        'end_sigma':0.304,
        'start_step_size':3.,
        'end_step_size':3.
    },
    {
        'layer':'loss1/classifier',
        'iter_n':5,
        'start_sigma':0.304,
        'end_sigma':0.2,
        'start_step_size':3.,
        'end_step_size':1.
    }
]

background_color = np.float32([85.0, 98.0, 112.0]) # "Slate" color -- tends to not work as well

for i in range(901,1000):
    imagenet_class = i
    gen_image = np.random.normal(background_color, 8, (original_w, original_h, 3))
    gen_image = deepdraw(net, gen_image, octaves, focus=imagenet_class, random_crop=True)
    img_fn = '_' .join([model_name, "deepdraw_denoised", str(imagenet_class)+'.png'])
    PIL.Image.fromarray(cv2.bilateralFilter(np.uint8(gen_image), -1, 25, 6)).save('./' + img_fn)
    # PIL.Image.fromarray(np.uint8(gen_image)).save('./' + img_fn)

```

In [ ]:

In [ ]: